All 26 model / reasoning combinations

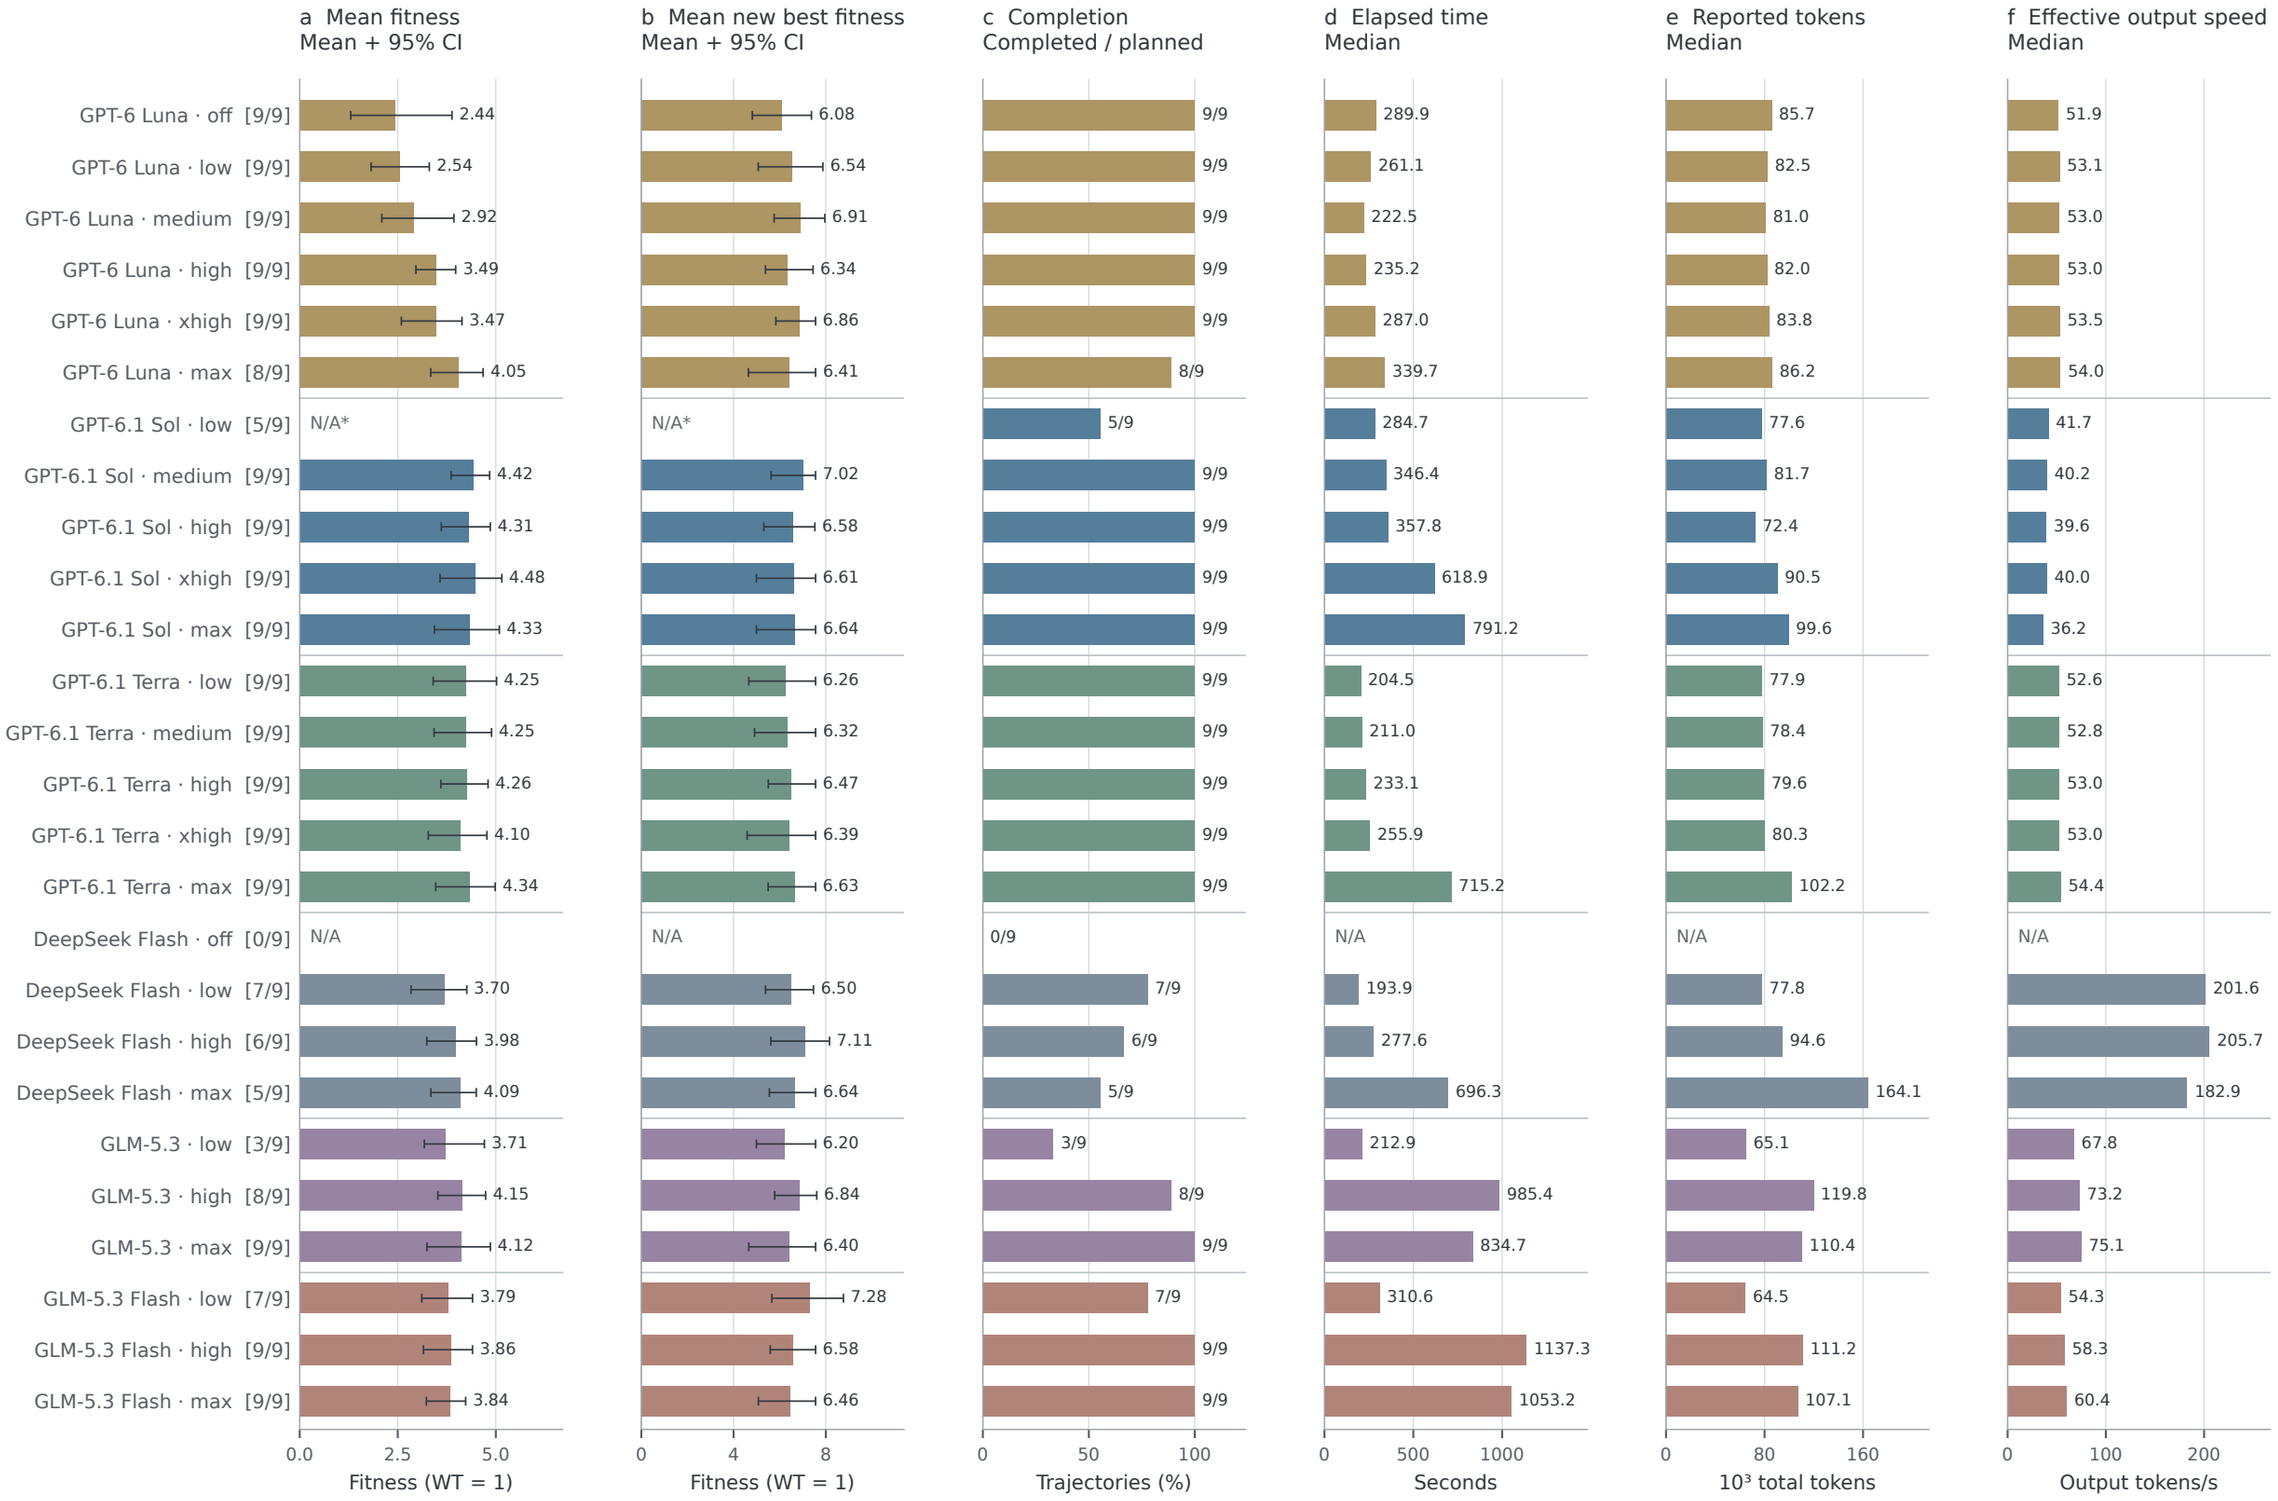

Phase 1: 3 data splits × 3 independent trajectories per combination. Labels show completed/planned trajectories.

Quality is conditional on completion; means give each split equal weight. Whiskers: audited 95% hierarchical bootstrap CI.

N/A: no completed trajectory. N/A*: at least one planned split has no completed trajectory; quality cannot be estimated.

Incomplete trajectories: 32. Missing usage: 11 calls (0 completed trajectories affected). Costs use completed trajectories only.

Speed = Σ output tokens / Σ full generation seconds per trajectory, then median; includes reasoning/waiting. Shared preparation excluded.